

15-213 Recitation: Data Lab

Jack Biggs

25 Jan 2016

Agenda

- Introduction
- Course Details
- Data Lab
 - Getting started
 - Running your code
 - ANSI C
- Bits & Bytes
- Integers
- Floating Point

Introduction

- Welcome to 15-213/18-213/15-513!
- Recitations are for...
 - Reviewing lectures
 - Discussing homework problems
 - Interactively exploring concepts
 - Previewing future lecture material
- Please, **please** ask questions!

Course Details

- How do I get help?
 - Course website: <http://cs.cmu.edu/~213>
 - Office hours: **5-9PM** from Sun-Thu in Wean 5207
 - Staff mailing list: 15-213-staff@cs.cmu.edu
 - *Definitely* consult the course textbook
 - **Carefully read the assignment writeups!**
- All labs are submitted on Autolab.
- All labs should be worked on using the **shark machines**.

Data Lab: Getting Started

- Download lab file (`datalab-handout.tar`)
 - Upload tar file to **shark** machine
 - `cd <my course directory>`
 - `tar xpvf datalab-handout.tar`
- `<filename>: Permission denied`
 - `chmod +x <filename>`
- Upload `bits.c` file to Autolab for submission

Data Lab: Running your code

- `dlc`: a modified C compiler that interprets *ANSI C only*
- `btest`: runs your solutions on random values
- `bddcheck`: exhaustively tests your solutions
 - Checks all values, formally verifying the solution
- `driver.pl`: Runs both `dlc` and `bddcheck`
 - Exactly matches Autolab's grading script
 - You will likely only need to submit once
- For more information, **read the writeup**
 - Available under assignment page as “**View writeup**”
 - **Read it. Read the writeup... please.**

Data Lab: What is ANSI C?

This is *not* ANSI C.

```
unsigned int foo(unsigned int x)
{
    x = x * 2;
    int y = 5;

    if (x > 5) {
        x = x * 3;
        int z = 4;
        x = x * z;
    }

    return x * y;
}
```

Within two braces, all *declarations* must go before any *expressions*.

Data Lab: What is ANSI C?

This is ANSI C.

```
unsigned int foo(unsigned int x)
{
    int y = 5;
    x = x * 2;

    if (x > 5) {
        int z = 4;
        x = x * 3;
        x = x * z;
    }

    return x * y;
}
```

This is *not* ANSI C.

```
unsigned int foo(unsigned int x)
{
    x = x * 2;
    int y = 5;

    if (x > 5) {
        x = x * 3;
        int z = 4;
        x = x * z;
    }

    return x * y;
}
```


Bits & Bytes: Unsigned integers

- An unsigned number represents positive numbers between 0 and 2^k-1 , where k is the numbers of bits used.
- Subtracting 1 from 0 will *underflow* to the highest value.
- Adding 1 to the highest value will *overflow* to 0

An 8-bit unsigned integer:

1	0	0	1	0	1	0	1
2^7			2^4		2^2		2^0


+ + + = 149

Bits & Bytes: Two's Complement

- In C, a *signed* number represents numbers between $[-2^{k-1}, 2^{k-1}-1]$, where k is the number of bits used.
- Overflow and underflow (from $\text{max} > 0$ and $\text{min} < 0$ values) is *undefined* with signed numbers in C
 - Depending on the underlying architecture, signed overflow / underflow could modulo, do nothing, or even abort the program
- The highest-level bit is set to 1 in negative numbers.
- To get the negative value of a positive number x , invert the bits of x and add 1.


Bits & Bytes: Two's Complement

From positive to negative:



Bits & Bytes: Two's Complement

From positive to negative:



Bits negated

Bits & Bytes: Two's Complement

From positive to negative:


0	0	0	1	0	1	0	1
1	1	1	0	1	0	1	0
1	1	1	0	1	0	1	1

Bits negated

Add one


Bits & Bytes: Two's Complement

From negative to positive:



Bits & Bytes: Two's Complement

From negative to positive:



Bits negated

Bits & Bytes: Two's Complement

From negative to positive:


1	1	0	1	0	1	0	1
0	0	1	0	1	0	1	0
0	0	1	0	1	0	1	1

Bits negated

Add one


Bits & Bytes: Two's Complement with 3 bits

- Why would anybody want to do this?
 - Uses the same circuitry for addition and subtraction!
- Note that there is no positive 4: the two's complement of -4 with three bits is -4




Bits & Bytes: Two's Complement with 3 bits

- Why would anybody want to do this?
 - **Uses the same circuitry for *addition* and subtraction!**
- Note that there is no positive 4: the two's complement of -4 with three bits is -4



Bits & Bytes: Two's Complement with 3 bits

- Why would anybody want to do this?
 - **Uses the same circuitry for addition and *subtraction*!**
- Note that there is no positive 4: the two's complement of -4 with three bits is -4




Bits & Bytes: Two's Complement with 3 bits

Negating The Minimum

$$-4 = 100$$

$$\sim(-4) = 011$$

$$\sim(-4) + 1 = 100$$



Bits & Bytes: Logical Operators

AND: &&

15 && 18 =

OR: ||

513 || 0 =

EQ: ==

15 == 18 =

NOT: !

!15213 =

Bits & Bytes: Logical Operators

AND: &&

15 && 18 = 1

OR: ||

513 || 0 =

EQ: ==

15 == 18 =

NOT: !

!15213 =

Bits & Bytes: Logical Operators

AND: &&

15 && 18 = 1

OR: ||

513 || 0 = 1

EQ: ==

15 == 18 =

NOT: !

!15213 =

Bits & Bytes: Logical Operators

AND: &&

15 && 18 = 1

OR: ||

513 || 0 = 1

EQ: ==

15 == 18 = 0

NOT: !

!15213 =

Bits & Bytes: Logical Operators

AND: &&

15 && 18 = 1

OR: ||

513 || 0 = 1

EQ: ==

15 == 18 = 0

NOT: !

!15213 = 0

Bits & Bytes: Bitwise Operators

AND: &

01100101

& 11101101

OR: |

01100101

| 11101101

XOR: ^

01100101

^ 11101101

NOT: ~

~11101101

Bits & Bytes: Bitwise Operators

AND: &

```
  01100101
& 01101101
-----
  01100101
```

OR: |

```
  01100101
| 01101101
-----
```

XOR: ^

```
  01100101
^ 01101101
-----
```

NOT: ~

```
~01101101
-----
```

Bits & Bytes: Bitwise Operators

AND: &

	01100101
&	11101101
	<hr/>
	01100101

OR: |

	01100101
	11101101
	<hr/>
	11101101

XOR: ^

	01100101
^	11101101
	<hr/>

NOT: ~

~	11101101
	<hr/>

Bits & Bytes: Bitwise Operators

AND: &

	01100101
&	11101101
	<hr/>
	01100101

OR: |

	01100101
	11101101
	<hr/>
	11101101

XOR: ^

	01100101
^	11101101
	<hr/>
	10001000

NOT: ~

~	11101101
	<hr/>

Bits & Bytes: Bitwise Operators

AND: &

	01100101
&	11101101
	<hr/>
	01100101

OR: |

	01100101
	11101101
	<hr/>
	11101101

XOR: ^


	01100101
^	11101101
	<hr/>
	10001000

NOT: ~


	01100101
~	11101101
	<hr/>
	00010010

Bits & Bytes: Shifting

Shifting modifies the positions of bits in a number:



Shifting right on a signed number will *extend the sign*:




(If the sign bit is zero, it will fill in with zeroes instead.)

This is known as “arithmetic” shifting.

Bits & Bytes: Shifting

Shifting right on an *unsigned* number will fill in with 0.



This is known as “logical” shifting.

Arithmetic shifting is useful for preserving the sign when dividing by a power of 2.

We get around this when we don't need it by using *bitmasks*.

In other languages, such as Java, it is possible to choose shifting operators, regardless of the type of integer. In C, however, it depends on the signedness.

Bits & Bytes: Endianness (Byte Order)


- Endianness describes which byte in a number comes first in memory. This is important for bomb lab.
- Little-Endian machines store the lowest-order byte first.
 - Intel machines (the shark machines!) are little-endian.

0xdeadbeef: 0x ef be ad de

- Big-Endian machines store the highest-order byte first.
 - The Internet is big-endian
 - How we think about binary numbers normally

0xdeadbeef: 0x de ad be ef

Floating Point: “Fixed” Point Representation



- Bits to the right of the “binary point” represent smaller fractions
- Difficult to represent a wide range of numbers
 - In this example, can't represent a number larger than 16
 - Can we sacrifice a bit of precision to accomplish this?

Floating Point: Scientific Notation

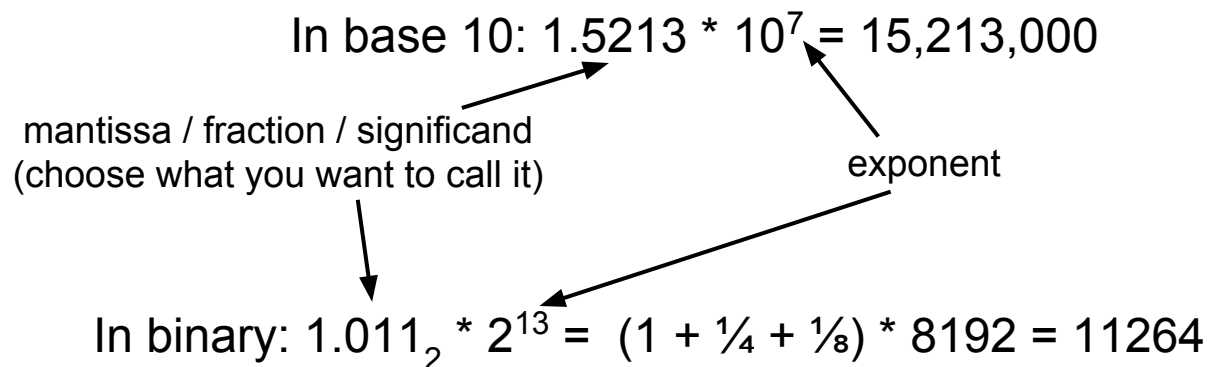
- In Scientific Notation, we represent a number as a fraction multiplied by an exponentiated scaling factor.

In base 10: $1.5213 * 10^7 = 15,213,000$


mantissa / fraction / significand
(choose what you want to call it)

exponent

In binary: $1.011_2 * 2^{13} = (1 + \frac{1}{4} + \frac{1}{8}) * 8192 = 11264$



Floating Point: IEEE Standard




Floating Point: Sign and Exponent



- If sign is 1, then the number is negative.
- The exponent determines three different value types.
 - Normalized: $0 \neq \text{exp} \neq 1111_2 \dots$
 - Mantissa = $s * 1.\text{frac}$
 - Denormalized: $\text{exp} = 0$
 - Mantissa = $s * 0.\text{frac}$
 - Special: $\text{exp} = 1111_2 \dots$
- ***Neither*** exp nor frac use two's complement!

Floating Point Example: Normalized


Consider a floating point implementation based on the IEEE floating point standard. This implementation omits the sign bit, and uses 4 bits for the exponent and 4 bits for the fraction.



$$E = \text{exp} - \text{bias}, \text{ where } \text{bias} = 2^{k-1} - 1.$$

Floating Point Example: Normalized


Consider a floating point implementation based on the IEEE floating point standard. This implementation omits the sign bit, and uses 4 bits for the exponent and 4 bits for the fraction.



$$E = \text{exp} - \text{bias}, \text{ where } \text{bias} = 2^{k-1} - 1.$$
$$\text{bias} = 2^{4-1} - 1 = 7.$$

Floating Point Example: Normalized

Consider a floating point implementation based on the IEEE floating point standard. This implementation omits the sign bit, and uses 4 bits for the exponent and 4 bits for the fraction.




$$E = \text{exp} - \text{bias}, \text{ where } \text{bias} = 2^{k-1} - 1.$$

$$\text{bias} = 2^{4-1} - 1 = 7.$$

$$E = 13 - 7 = 6$$

Floating Point Example: Normalized

Consider a floating point implementation based on the IEEE floating point standard. This implementation omits the sign bit, and uses 4 bits for the exponent and 4 bits for the fraction.



$$E = \text{exp} - \text{bias}, \text{ where } \text{bias} = 2^{k-1} - 1.$$


$$\text{bias} = 2^{4-1} - 1 = 7.$$

$$E = 13 - 7 = 6$$

Fraction has an implied leading 1.

Floating Point Example: Normalized

Consider a floating point implementation based on the IEEE floating point standard. This implementation omits the sign bit, and uses 4 bits for the exponent and 4 bits for the fraction.



$$E = \text{exp} - \text{bias}, \text{ where bias} = 2^{k-1} - 1.$$

$$\text{bias} = 2^{4-1} - 1 = 7.$$


$$E = 13 - 7 = 6$$

Fraction has an implied leading 1.

$$\text{Mantissa} = 1.0110_2$$

Floating Point Example: Normalized

Consider a floating point implementation based on the IEEE floating point standard. This implementation omits the sign bit, and uses 4 bits for the exponent and 4 bits for the fraction.



$$E = \text{exp} - \text{bias}, \text{ where bias} = 2^{k-1} - 1.$$

$$\text{bias} = 2^{4-1} - 1 = 7.$$


$$E = 13 - 7 = 6$$

Fraction has an implied leading 1.

$$\text{Mantissa} = 1.0110_2$$


$$2^6 * 1.0110_2 = 64 * (1 + \frac{1}{4} + \frac{1}{8}) = \mathbf{88}$$

Floating Point Example: Denormalized




$E = 1 - \text{bias}$, and $\text{bias} = 7$ as before.

Floating Point Example: Denormalized



$E = 1 - \text{bias}$, and $\text{bias} = 7$ as before.
Fraction has an implied leading 0.

Floating Point Example: Denormalized




$E = 1 - \text{bias}$, and $\text{bias} = 7$ as before.

Fraction has an implied leading 0.

Fraction is equal to 0.1100_2

Floating Point Example: Denormalized




$E = 1 - \text{bias}$, and $\text{bias} = 7$ as before.

Fraction has an implied leading 0.


Fraction is equal to 0.1100_2

Final answer: $2^{-6} * (0 + \frac{1}{2} + \frac{1}{4}) = \mathbf{0.01171875}$

Floating Point Example: Special Values




- Exp is all 1
- If fraction is all 0, then represents infinity
 - Also, -Infinity (if we had a sign bit)




- If fraction $\neq 0$, then represents NaN (Not a Number!)
- Sign bit doesn't *really* matter, but either can turn up
 - (Mostly from division errors)

Floating Point Example: Limits



- What is the largest denormalized number?
- What is the smallest normalized number?
- What is the largest finite number it can represent?
- What is the smallest non-zero value it can represent?

Floating Point Example: Limits




- **What is the largest denormalized number?**
- What is the smallest normalized number?
- What is the largest finite number it can represent?
- What is the smallest non-zero value it can represent?

$$0000\ 1111 = 0.1111_2 * 2^{-6} = 0.0146484375$$

(recall that $E = 1 - \text{bias}$, and $\text{bias} = 7$ in this example)

Floating Point Example: Limits




- What is the largest denormalized number?
- **What is the smallest normalized number?**
- What is the largest finite number it can represent?
- What is the smallest non-zero value it can represent?

0001 0000

$$E = 1 - 7 = -6$$

$$\text{Answer: } 1.0000_2 * 2^{-6} = 2^{-6} = 1/64$$

Floating Point Example: Limits




- What is the largest denormalized number?
- What is the smallest normalized number?
- **What is the largest finite number it can represent?**
- What is the smallest non-zero value it can represent?

1110 1111

$$E = 14 - 7 = 7$$

$$\text{Answer: } 1.1111_2 * 2^7 = 248$$

Floating Point Example: Limits



- What is the largest denormalized number?
- What is the smallest normalized number?
- What is the largest finite number it can represent?
- **What is the smallest non-zero value it can represent?**

0000 0001

$$= 0.0001_2 * 2^{-6} = 0.0009765625$$

(recall that $E = 1 - \text{bias}$, and $\text{bias} = 7$ in this example)

Floating Point: Rounding

1.**B****B****G****R****X****X****X**

*In the below examples,
imagine the underlined part
as a fraction.*

- **Guard Bit**: the least significant bit of the resulting number
- **Round Bit**: the first bit removed from rounding
- **Sticky Bits**: all bits after the round bit, OR'd together

Examples of rounding cases, including rounding to nearest even number

- 1.**10****|****11**: More than $\frac{1}{2}$, round up: 1.**11**
- 1.**10****|****10**: Equal to $\frac{1}{2}$, round down *to even*: 1.**10**
- 1.**01****|****01**: Less than $\frac{1}{2}$, round down: 1.**01**
- 1.**01****|****10**: Equal to $\frac{1}{2}$, round up *to even*: 1.**10**
- 1.**01****|****00**: Equal to 0, do nothing: 1.**01**
- 1.**00****|****00**: Equal to 0, do nothing: 1.**00**

All other cases involve either rounding up or down - *try them!*

Questions?

- Remember, data lab is due this Thursday!
 - You really should have started already!
- Read the lab writeup.
 - **Read the lab writeup.**
 - ***Read the lab writeup.***
 - Please. :)