

15-213 Recitation 5: Attack Lab

8 Feb 2016

Ralf Brown and the 15-213 staff

Agenda

- Reminders
- Stacks
- Attack Lab Overview
- Appendix: Arrays
- Appendix: Structs
- Appendix: More Assembly



Reminders

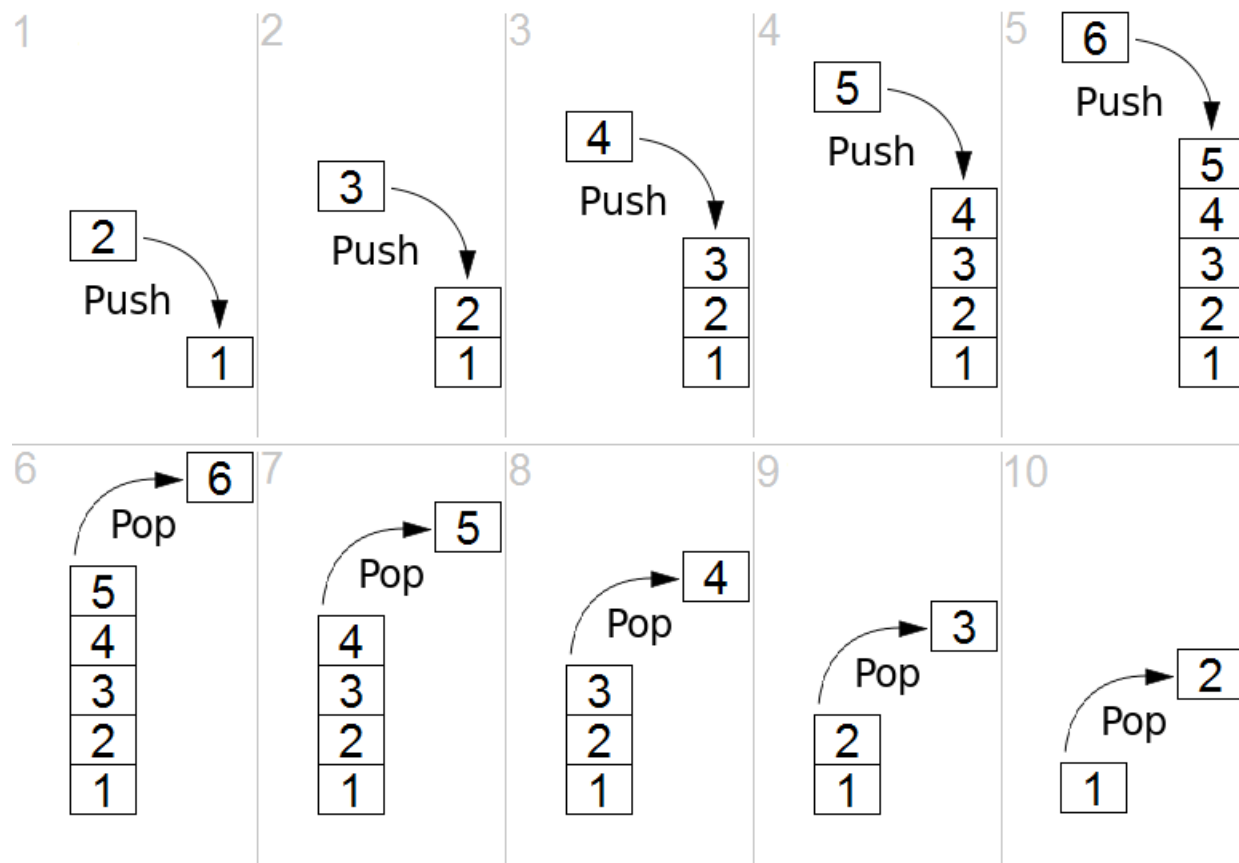
- Bomb lab is due **tomorrow!**
 - “But if you wait until the last minute, it only takes a minute!” - ***NOT!***
 - Don't waste your grace days on this assignment
- Attack lab will be released **tomorrow!**



Image credit: tzunghaor / openclipart.org

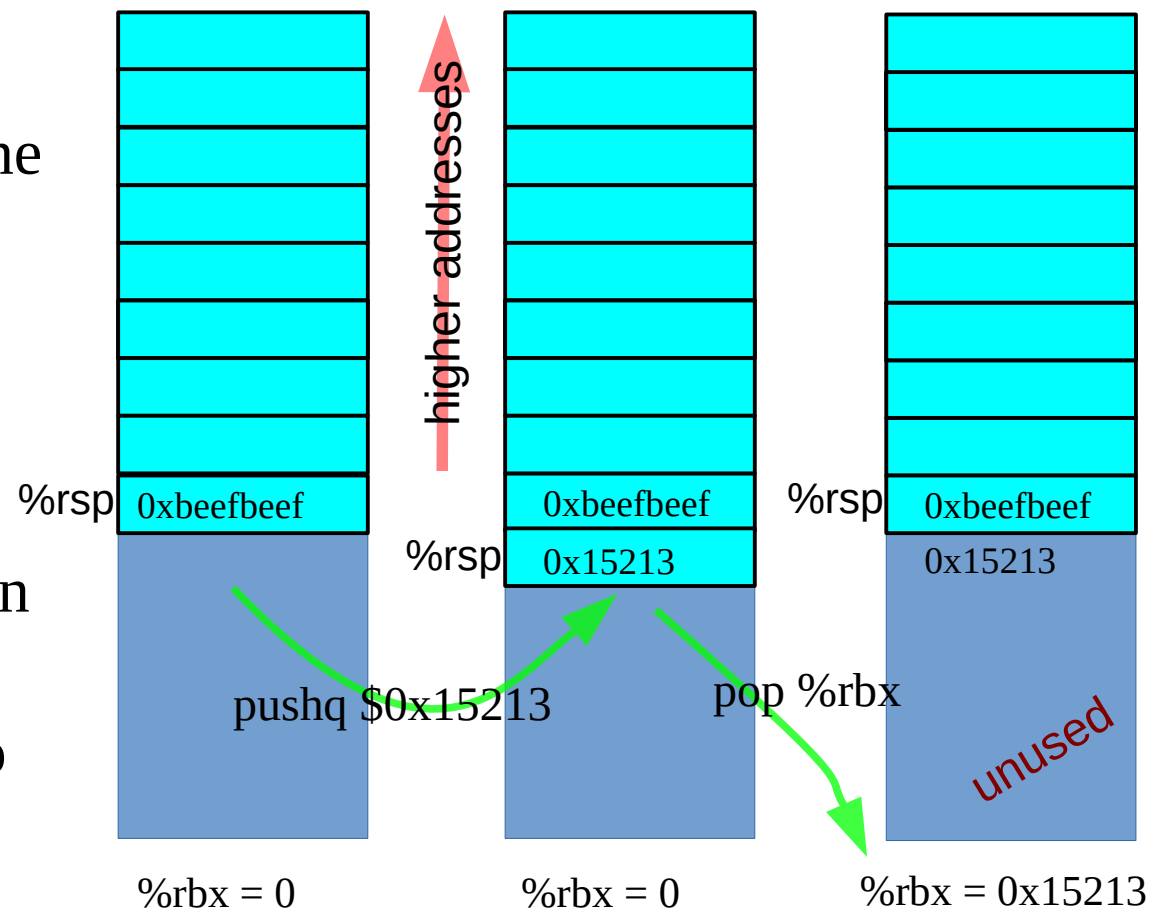
Stacks

- Last-In, First-Out
 - just like a stack of plates
 - pushes and pops to preserve registers must be in **opposite** order
- x86 stack grows **down**
 - lowest address is “top”



Stacks

- `%rsp` contains the address of the topmost element of the stack
- `pushq {value}` is same as
`sub $8,%rsp`
`mov {value}, (%rsp)`
- only constants and registers can be pushed
- `popq {reg}` is equivalent to
`mov (%rsp), {reg}`
`add $8, %rsp`



Stacks

- Stacks are useful for recursion
 - each call of the function gets a separate copy of arguments and local variables
 - the separate copy is needed for a limited time only – until the call returns
 - callee always ends before caller
- This is called “stack discipline”
- Stack space is allocated in **Frames**
 - state for a single instantiation of a function

Register Saving Conventions

■ Caller-Saved

- called function may do as it wishes with the register
- must save/restore register in caller's stack frame if it still needs the value after a function call
- registers used as function arguments are always caller-saved
- result register `%rax` is also caller-saved

■ Callee-Saved

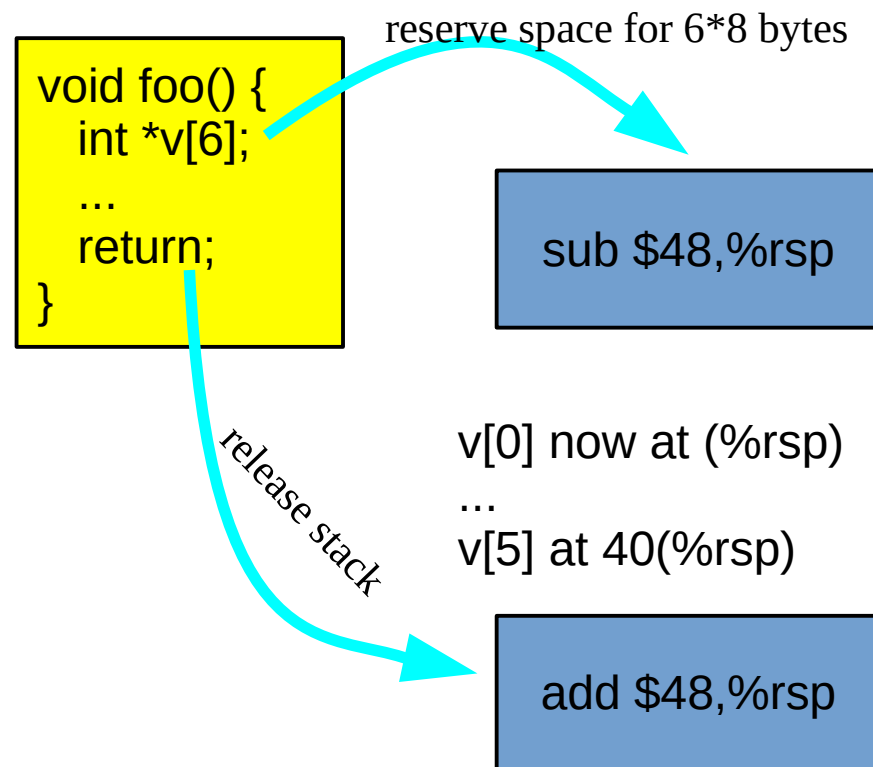
- if the function wants to change the register, it must save the original value in its stack frame and restore it before returning
- the calling function may store temporary values across function calls in callee-saved registers

x86-64 Register Usage Conventions

%rax	return value	%r8	argument #5
%rbx	callee saves	%r9	argument #6
%rcx	argument #4	%r10	caller saves
%rdx	argument #3	%r11	caller saves
%rsi	argument #2	%r12	callee saves
%rdi	argument #1	%r13	callee saves
%rsp	stack pointer	%r14	callee saves
%rbp	callee saves	%r15	callee saves

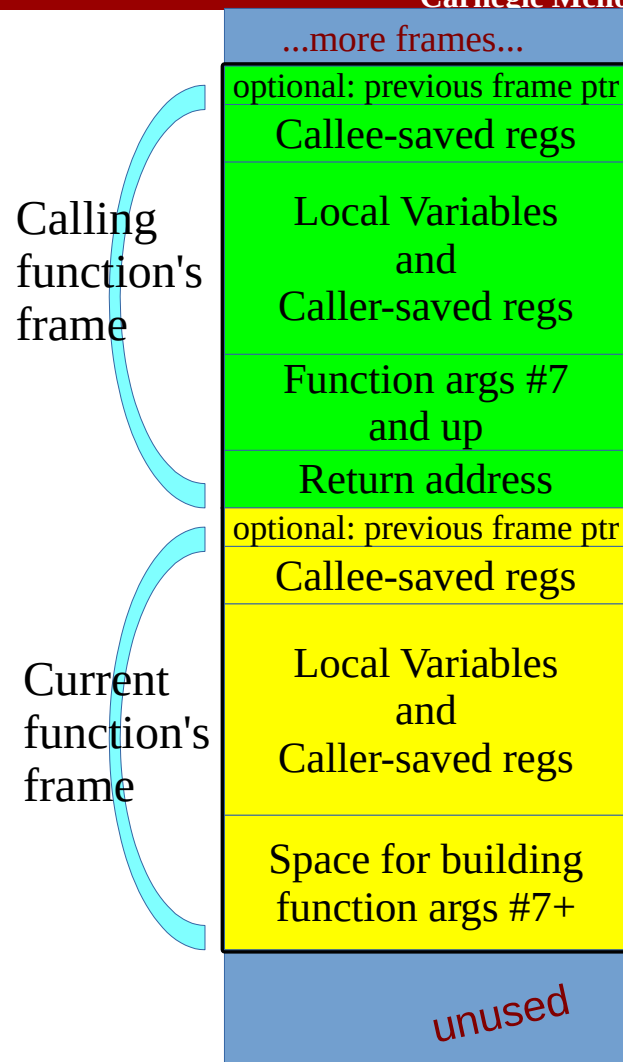
Local Variables

- Local variables which can't be stored in registers are stored on the stack
 - this includes arrays, structs, and anything which has its address taken
- Storage is allocated by simply decrementing `%rsp` the appropriate amount
- Cleanup consists of incrementing `%rsp` to free the stack space



Stack Frames

- A frame can have many parts, but only those needed by the function are actually present
 - we consider the function args and return address to be part of the *caller's* frame because they are pushed *before* control transfers to the callee
- Function args 7+ are stored in order from lowest address to highest
 - equivalent to pushing in reverse order
 - access by offsetting from own stack frame
- When present, frame pointers let you traverse the chain of stack frames



Smashing the Stack

- What if a function has a bug that causes it to write beyond its own frame, possibly on some specific program input?
 - We call such overwrites “smashing the stack”
 - If the return address is overwritten, we could end up *anywhere* when the function returns
 - Classic article: “[Smashing the Stack for Fun and Profit](#)” by Aleph One (1996).

Attack Lab

- The goal of this lab is to figure out some input to be fed to a program to make it do things it was never designed for
 - You'll do this by smashing the stack
- In the process, you will learn how to prevent and defend against such attacks
- **READ THE WRITEUP!** It shows you the techniques and helper programs you'll need to successfully exploit your target program.


Attack Lab: Code Injection

- You'll be given two compiled programs with a buffer-overflow vulnerability
- Part 1 has you smashing the stack to place code of your own on the stack to create the necessary conditions to call the target function
 - this means you'll have to craft the appropriate bytes in your attack string and then jump to them somehow once they're on the stack

Code Injection Example

- `strcpy` copies from `input` up to and including a NUL byte (`'\0'`)
- if there are more than 31 bytes before the NUL, `strcpy` will write beyond the end of `buf`
- the next thing on the stack is the return address for `vuln()`....


```
void vuln(char *input) {  
    char buf[32];  
    ...  
    strcpy(buf, input);  
    return;  
}
```



Code Injection Example

- Simplest case: we just need to jump to a different address
 - fill the buffer with anything, then overwrite the return address
 - if we need to jump to 0xFFF15213 instead of the original 0xFEEDBEEF, a possible value for input is
0x41 (32 times) 0x13 0x52 0xF1 0xFF 0x00
 - remember byte order!
- After the strcpy, the stack will look as shown


```
void vuln(char *input) {  
    char buf[32];  
    ...  
    strcpy(buf, input);  
    return;  
}
```



Code Injection Example

- If we need to insert actual code on the stack, we must first figure out the value of `%rsp`!
 - run a copy of the program and break at the call to `strcpy`
- Put your exploit code in the buffer, and overwrite the return address with the address of the buffer
 - use **gcc** and **objdump** to generate the byte sequences for your injected code
 - code can't contain NUL bytes (newlines if exploiting `gets`)

```
void vuln(char *input) {  
    char buf[32];  
    ...  
    strcpy(buf, input);  
    return;  
}
```



Attack Lab: The Stack Has Protections!

- Modern CPUs allow data areas to be marked as non-executable
 - so you can't place exploit code on the stack and jump to it
- OSes use Address Space Layout Randomization to keep addresses unpredictable
 - an attack that works on one run might not on the next!
- Compiled code commonly uses “stack canaries”
 - unpredictable values stored between on-stack buffers and return address
 - if the value changes, the program is aborted rather than returning from the function (because the return address might have been corrupted)

Attack Lab: Return-Oriented Programming

- In Part 2 of Attack Lab, the program has stack protections in place
- If we can't build our code on the stack, we have to find snippets of existing code (“gadgets”) that we can stitch together with an appropriate sequence of return addresses on the stack
 - hence the name ROP
 - we just need to find the right byte sequences – those exact instruction sequences don't actually have to be a deliberate part of the program!

ROP: Finding Gadgets

- Look for byte sequences corresponding to an interesting instruction followed by `ret`
 - e.g. `0x59 0xC3` would be `pop %rbx; ret`
- Need to be creative in the instructions we execute
 - we probably can't find `mov $0x15213, %rax`
 - but we could stitch together the equivalent from
 - `pop %rbx`
 - `mov %rbx, %rax`
 - and a value of `0x15213` on the stack

ROP: The Exploit

- Find the gadgets
- Overflow the buffer to overwrite the return address and higher stack addresses with ROP instructions
 - each is a gadget address followed by any data the gadget pops off the stack

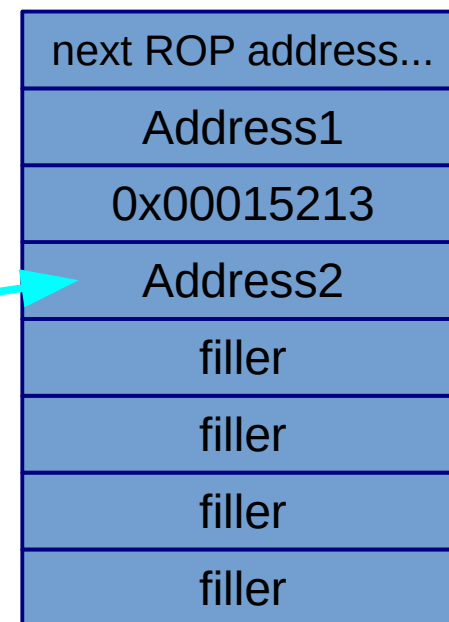
```
void vuln(char *input) {  
    char buf[32];  
    ...  
    strcpy(buf, input);  
    return;  
}
```

Address1:
mov %rbx, %rax
ret

Address2:
pop %rbx
ret

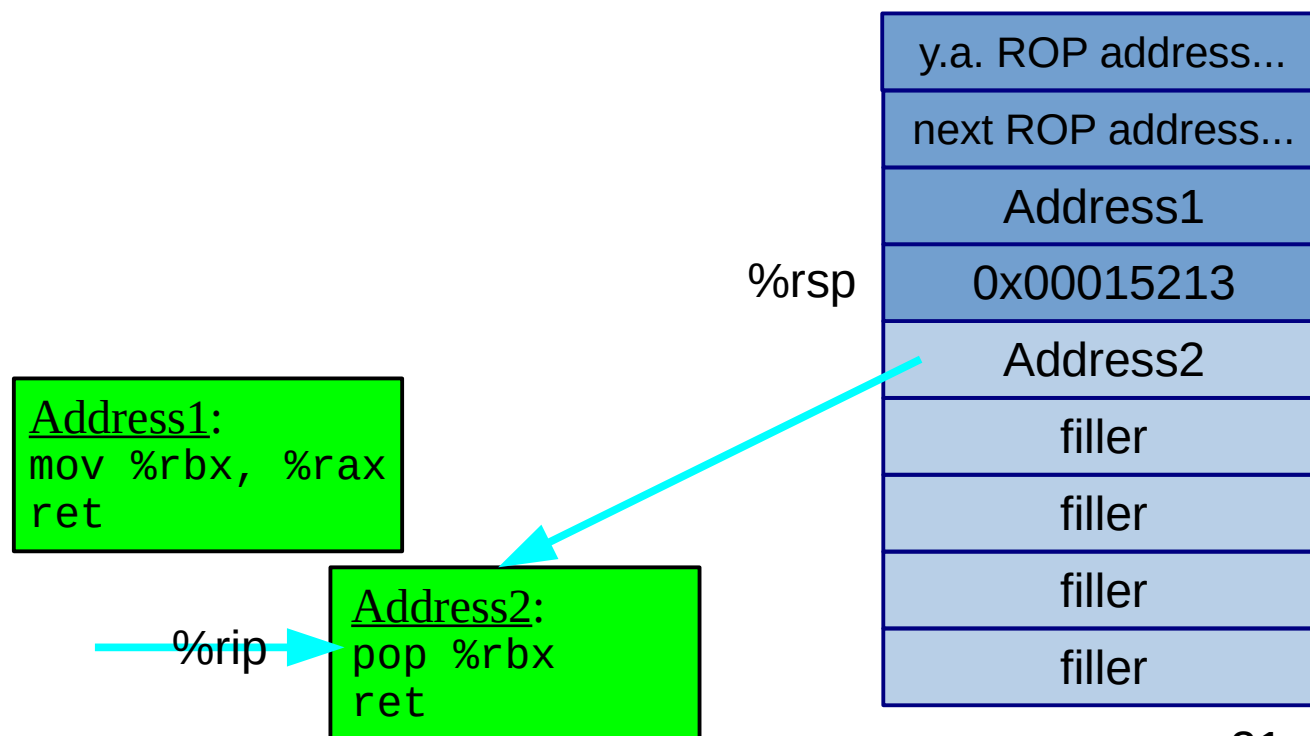
old return
address
was here

buf



ROP: The Exploit Runs

- When `vuln()` returns, we jump to Address2 instead of the original caller
 - `%rax = ?`
 - `%rbx = ?`




ROP: The Exploit Runs

- The gadget at Address2 pops the top item off the stack
 - `%rax = ?`
 - `%rbx = 0x00015213`

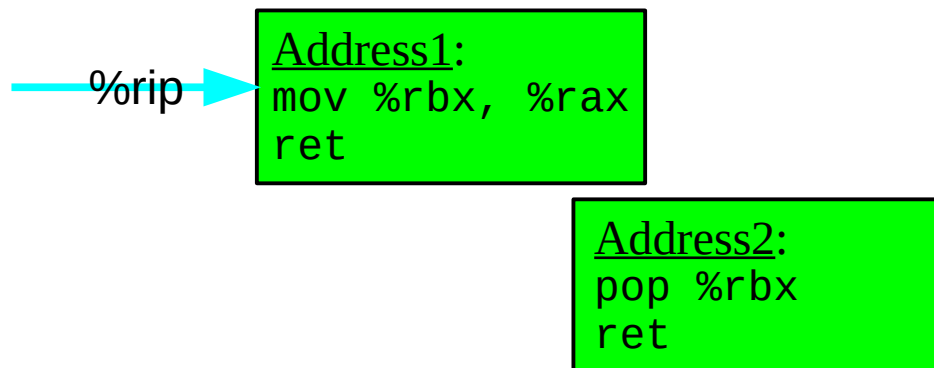
Address1:
`mov %rbx, %rax`
`ret`

→ `%rip` → Address2:
`pop %rbx`
`ret`




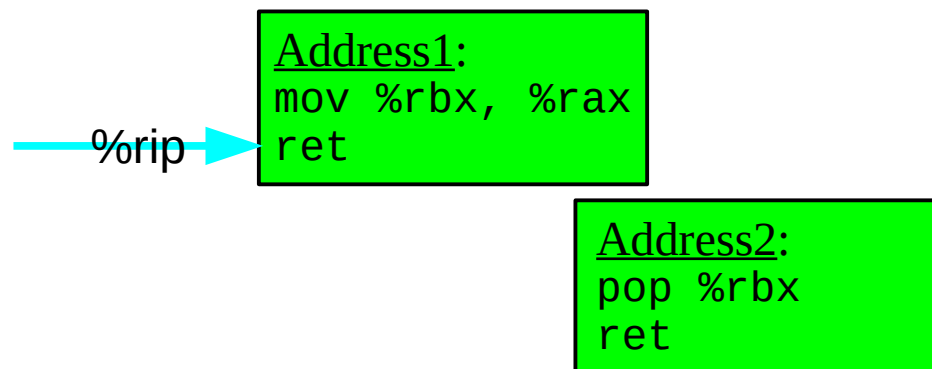
ROP: The Exploit Runs

- Next, the gadget at Address2 returns, which puts us at Address1
 - `%rax = ?`
 - `%rbx = 0x00015213`



ROP: The Exploit Runs

- The gadget at Address1 now copies %rbx into %rax
 - %rax = 0x00015213
 - %rbx = 0x00015213



ROP: The Exploit Runs

- Finally, the gadget at Address1 returns, taking us to the next gadget
 - `%rax = 0x00015213`
 - `%rbx = 0x00015213`
- We've now executed the equivalent of
 - `mov $0x15213, %rax`
- (with a side effect)

```
Address1:  
mov %rbx, %rax  
ret
```

```
Address2:  
pop %rbx  
ret
```



Attack Lab Tools

- **gcc -c file.s**

- convert the assembly code in file.s to object code in file.o

- **objdump -d file.o**

- disassemble the code in file.o; shows the actual bytes for the instructions

- **./hex2raw**

- convert hex codes into raw ASCII strings to pass to targets

- **gdb**

- determine stack addresses

- **paper and pencil**

- for drawing stack diagrams

More Useful GDB Commands

<code>x/[n]i <address></code>	disassemble n instructions at <code><address></code>
<code>b <loc> if <cond></code>	conditional breakpoint, stop only if <code><cond></code> true
<code>cond <bp> <cond></code>	add condition to existing breakpoint <code><bp></code>
<code><u>commands</u> <bp></code>	execute commands when breakpoint <code><bp></code> hit
<code><u>tbreak</u> <loc></code>	set temporary breakpoint – auto-deletes when hit!
<code><u>finish</u></code>	run until current frame (function) returns, and print return value
<code><u>layout</u> <u>asm</u></code>	split the screen into separate disassembly and command windows
<code><u>layout</u> <u>reg</u></code>	show register window as well (after <code>layout asm</code>)

If You Get Stuck

- Please read the writeup. *Please read the writeup.* Please read the writeup. ***Please read the writeup!***
- CS:APP Chapter 3
- View lecture notes and course FAQ at <http://www.cs.cmu.edu/~213>
- Office hours Sunday through Thursday 5:00-9:00pm in WeH 5207
- Post a **private** question on Piazza
- `man gdb`, `gdb's help` command

Remember...



Appendix: Arrays


- In C, the name of an array is interpreted as a pointer to the first element
 - A is the same as $\&A[0]$
- Array subscripting is just a synonym for pointer arithmetic:
 - $A[1]$ equals $*(A + 1)$
- This translates almost directly into assembly. $x = A[5]$ becomes

```
mov $5, %rax
mov {address of A}, %rbx
mov (%rbx, %rax, 4), %rdx
```
- We simply scale the index by the size of an element and add that to the starting address

Two-dimensional Arrays

- Arrays elements can themselves be arrays
- As with one-dimensional arrays, the elements are stored in order in memory
- C only supports compile-time sized multi-dimensional arrays – you need to compute the corresponding index as if the array were one-dimensional for run-time sizing

```
int A[3][2] = {  
    { 0, 1 },  
    { 10, 11 },  
    { 20, 21 } };
```



$$\&A[1][0] = A + 1*2 + 0 = A + 2$$

$$\&A[2][1] = A + 2*2 + 1 = A + 5$$

for $M \times N$ array, $\&A[i][j] = A + i*N + j$

Appendix: Structs

- Structures are a way to bundle together related data/variables
- Elements of a structure may be of any type, including pointers, arrays and *other* structures (no recursive regress allowed!)
- Structs can be the elements of an array
- Access parts of the structure by name, rather than by index as for arrays

```
struct info {  
    int whole_num;  
    double float_num;  
    char string[9];  
} S[5];  
/* init 2nd array element */  
S[1].whole_num = 15213;  
S[1].float_num = 15.213;  
strcpy(S[1].string, "15-213");
```


Structs

- Structures are a way to bundle together related data/variables
- Elements of a struct can be any type, including pointers, arrays and *other* structures (no recursive regress allowed!) **declared variable**
- Structs can be the elements of an array
- Access parts of the structure by name, rather than by index as for arrays

```
struct info {  
    int whole_num;  
    double float_num;  
    char string[9];  
} S[5];  
/* init 2nd array element */  
S[1].whole_num = 15213;  
S[1].float_num = 15.213;  
strcpy(S[1].string, "15-213");
```

type name

member

Structs

- The type name of a struct can be used in future variable declarations – given the declaration on the previous slide, we can now say

```
struct info T;
```

- This declares variable T to be of the same type as the elements of array S.


- The type name is optional when declaring a variable – that creates an *anonymous* structure type.
- You can also declare a structure type without declaring any variables:

```
struct st { int m1; float m2; };
```


Structs: Memory Layout

- Struct members are placed a multiple of their own size from the start of the struct
 - some architectures require such alignment to even access the member
 - x86 doesn't care, but will be slower whenever misalignment causes two memory accesses
- Minimize size by putting largest members first
 - array elements count as individual members

```
struct s1 {  
    int i;  
    double *p;  
    char s[9];  
};
```





```
struct s2 {  
    double *p;  
    int i;  
    char s[9];  
};
```



Arrays of Structs

- Requiring struct sizes to be a multiple of their strictest alignment supports arrays
- Otherwise, not all members would be aligned in every array element



Misaligned!

Appendix: More Assembly

- Some instructions you may encounter
 - `c1tq` (“Convert Long To Quad”) -- sign-extend `%eax` into `%rax`
 - `cmovX` (“Conditional Move”) -- executes move only if condition *X* is true
 - `movzbl` (“MOVE w/ Zero-extension, Byte to Long”)
 - `nop` (“No Operation”) -- do nothing
 - `nopl` (“No Operation, Long”) -- multi-byte instruction that does nothing
 - used to align function start addresses to a multiple of 16 bytes
 - `repz retq` – see Recitation 4
 - `mov %fs:0x28, %rax`
 - beyond the scope of the course; `%fs` is a *segment register*, which here is used to implement thread-local storage