

15-213 Recitation 14: Threads and Synchronization

18 April 2016

Ralf Brown and the 15-213 staff

Agenda

- Reminders
- Threads Revisited
- Synchronization

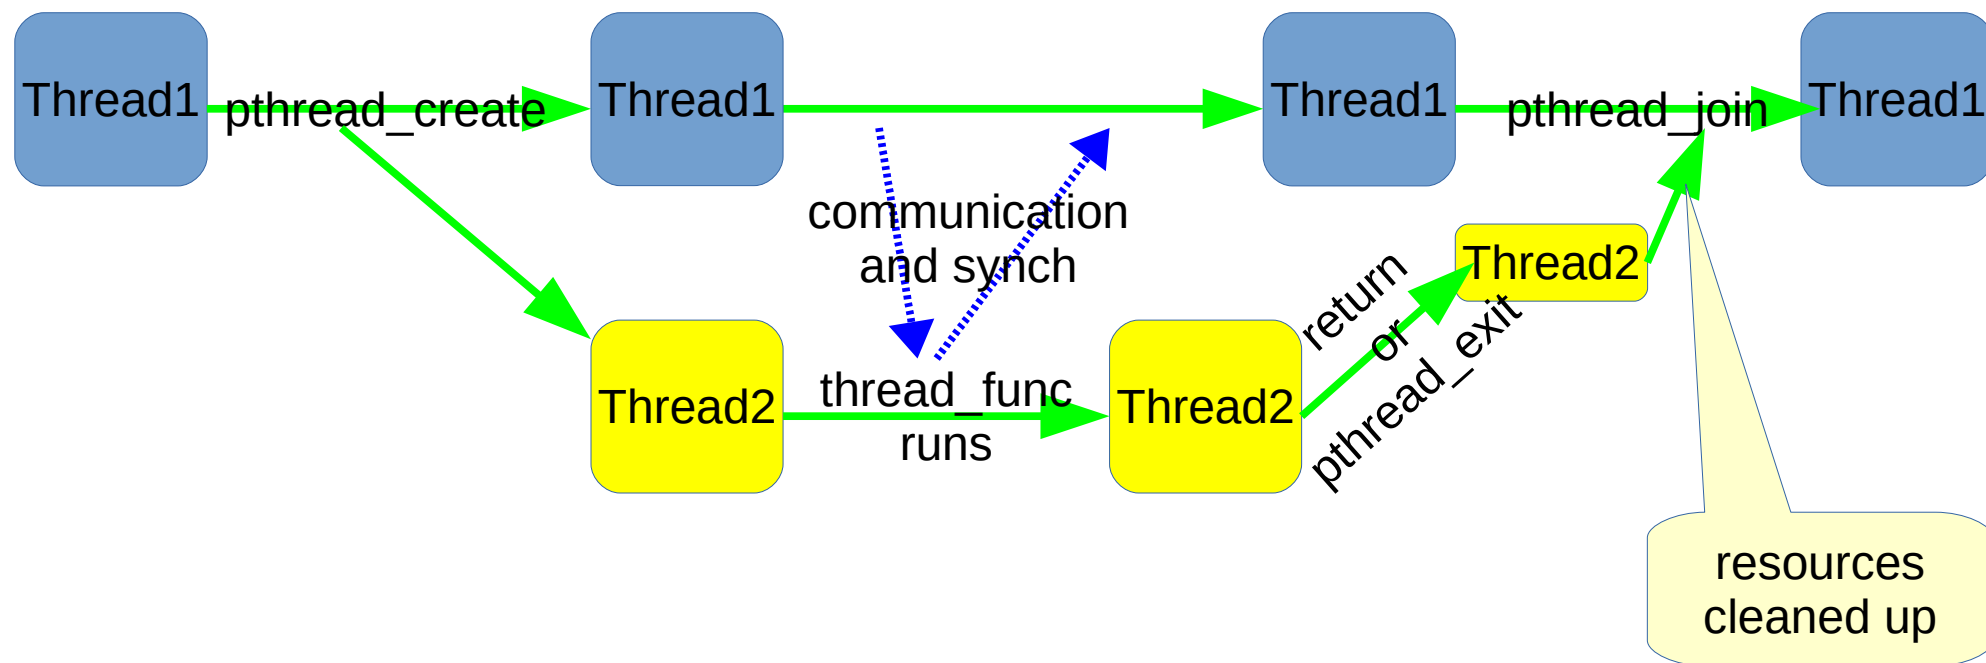


These aren't the only threads that can get tangled....

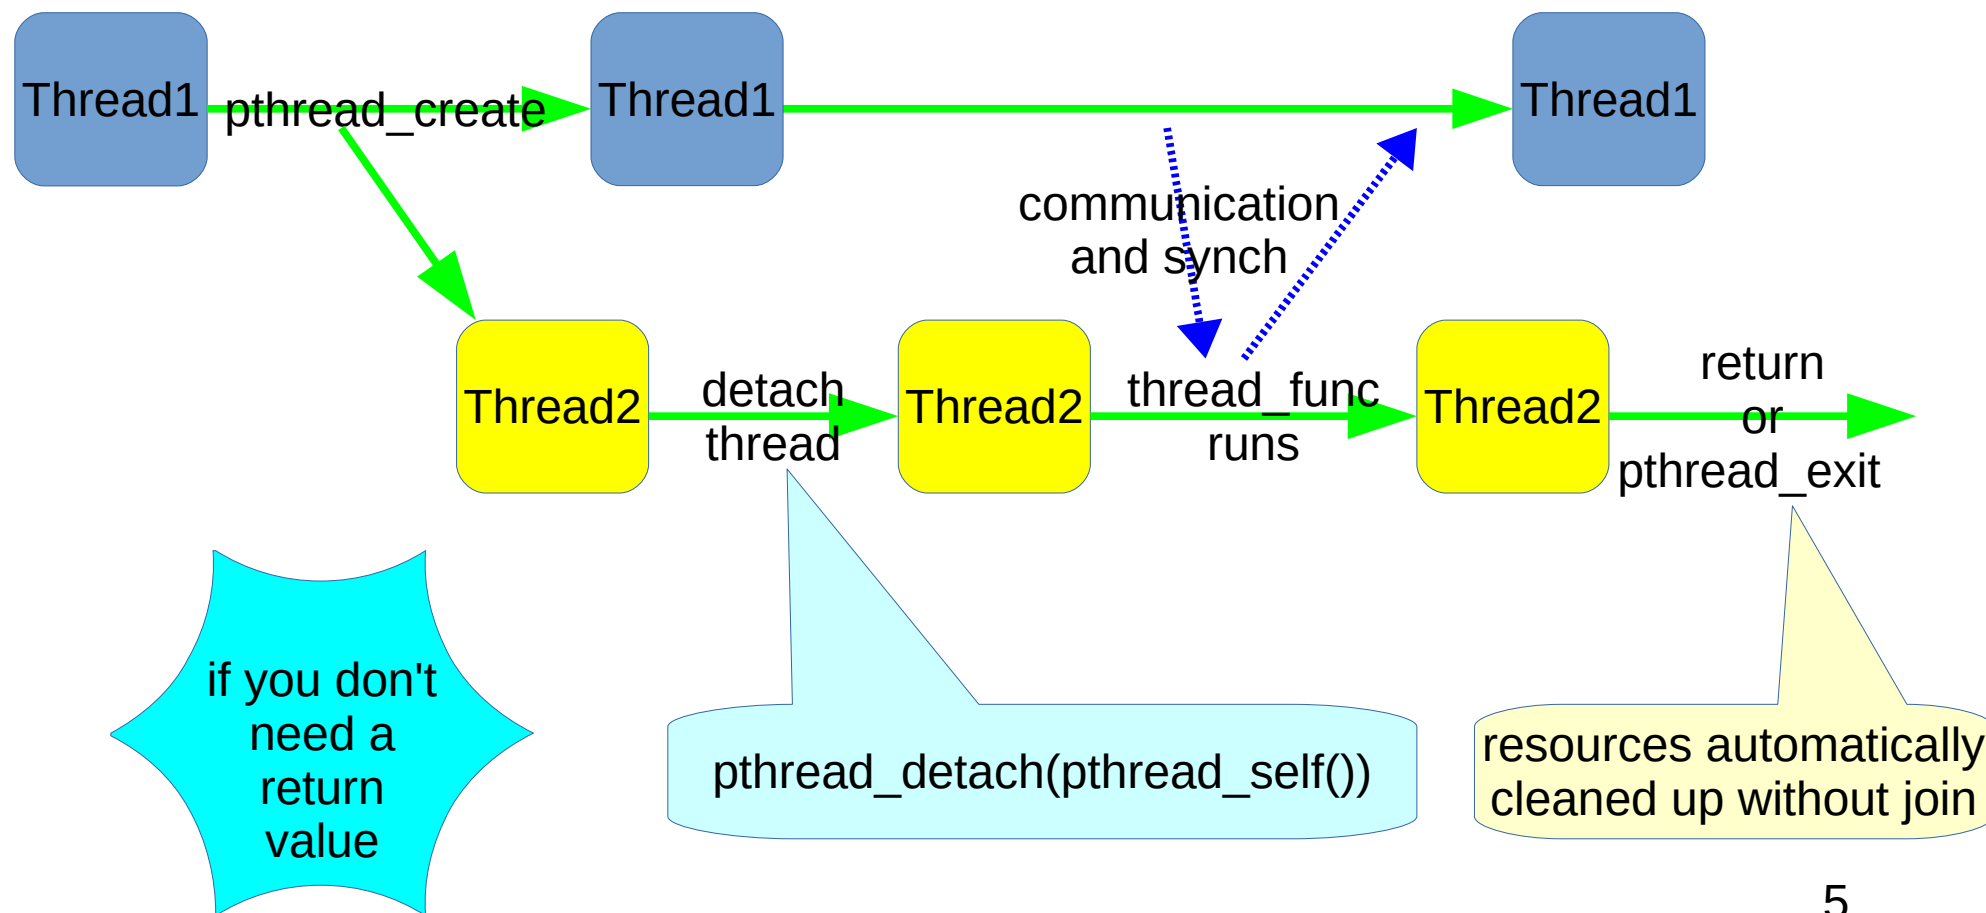
Reminders

- Proxylab is due a week from tomorrow
 - there are no grace days!
 - make your code robust against unexpected inputs

Thread Life Cycle



Thread Life Cycle



Using Threads

- Let's sum the elements of an array

Single-threaded

```
int *nums;

int sum_array(size_t n) {
    int sum = 0;
    // iterate over the elements of the array
    for (int i = 0; i < n; i++)
        sum += nums[i];
    return sum;
}
```

Using Threads: Summing an Array

Multi-threaded

```
// the main thread function - sum a section of the array
void *thread_fun(void *vargp) {
    int myid = *((int*)vargp);
    size_t start = myid * nelems_per_thread;
    size_t end = start + nelems_per_thread;
    size_t i;
    int sum = 0;
    for (i = start; i < end; i++)
        sum += nums[i];
    psum[myid] = sum;
    return NULL;
}
```

note: we're
omitting a
bunch of
variable
declarations

Using Threads: Summing an Array

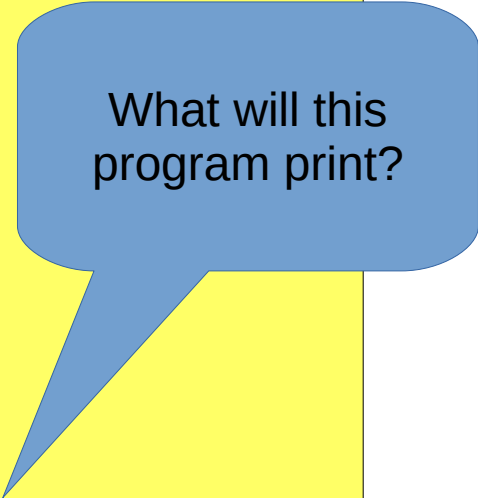
```
int sum_array(int nelems) {  
    int sum = 0;  
    // figure out how big the sections should be  
    nelems_per_thread = nelems / nthreads;  
    // create threads  
    for (i = 0; i < nthreads; i++) {  
        myid[i] = i;  
        Pthread_create(&tid[i], NULL, thread_fun, &myid[i]);  
    }  
    // wait for the threads to finish  
    for (i = 0; i < nthreads; i++)  
        Pthread_join(tid[i], NULL);  
    // collect results  
    for (i = 0; i < nthreads; i++)  
        sum += psum[i];  
    // add leftover elements  
    for (e = nthreads * nelems_per_thread; e < nelems; e++)  
        sum += nums[e];  
    return sum;  
}
```

why use an
array for myid?

Critical Sections and Shared Variables

```
volatile int total = 0;
void *incr(void *ptr) {
    for (int i = 0; i < *((int*)ptr); i++)
        total++;
    return NULL;
}

int main() {
    pthread_t tids[NTHREADS];
    int y = NINCR;
    for (int i = 0; i < NTHREADS; i++)
        Pthread_create(&tids[i], NULL, incr, &y);
    for (int i = 0; i < NTHREADS; i++)
        Pthread_join(tids[i], NULL);
    printf("total is: %d\n", total);
    return 0;
}
```



What will this program print?

Critical Sections and Shared Variables

```
volatile int total = 0;
void *incr(void *ptr) {
    for (int i = 0; i < *((int*)ptr); i++)
        total++;
    return NULL;
}

int main() {
    pthread_t tids[NTHREADS];
    int y = NINCR;
    for (int i = 0; i < NTHREADS; i++)
        Pthread_create(&tids[i], NULL, incr, &y);
    for (int i = 0; i < NTHREADS; i++)
        Pthread_join(tids[i], NULL);
    printf("total is: %d\n", total);
    return 0;
}
```

We have NTHREADS
incrementing the
total NINCR times
each

but the total may
be much less than
 $NINCR * NTHREADS$!

Critical Sections and Shared Variables

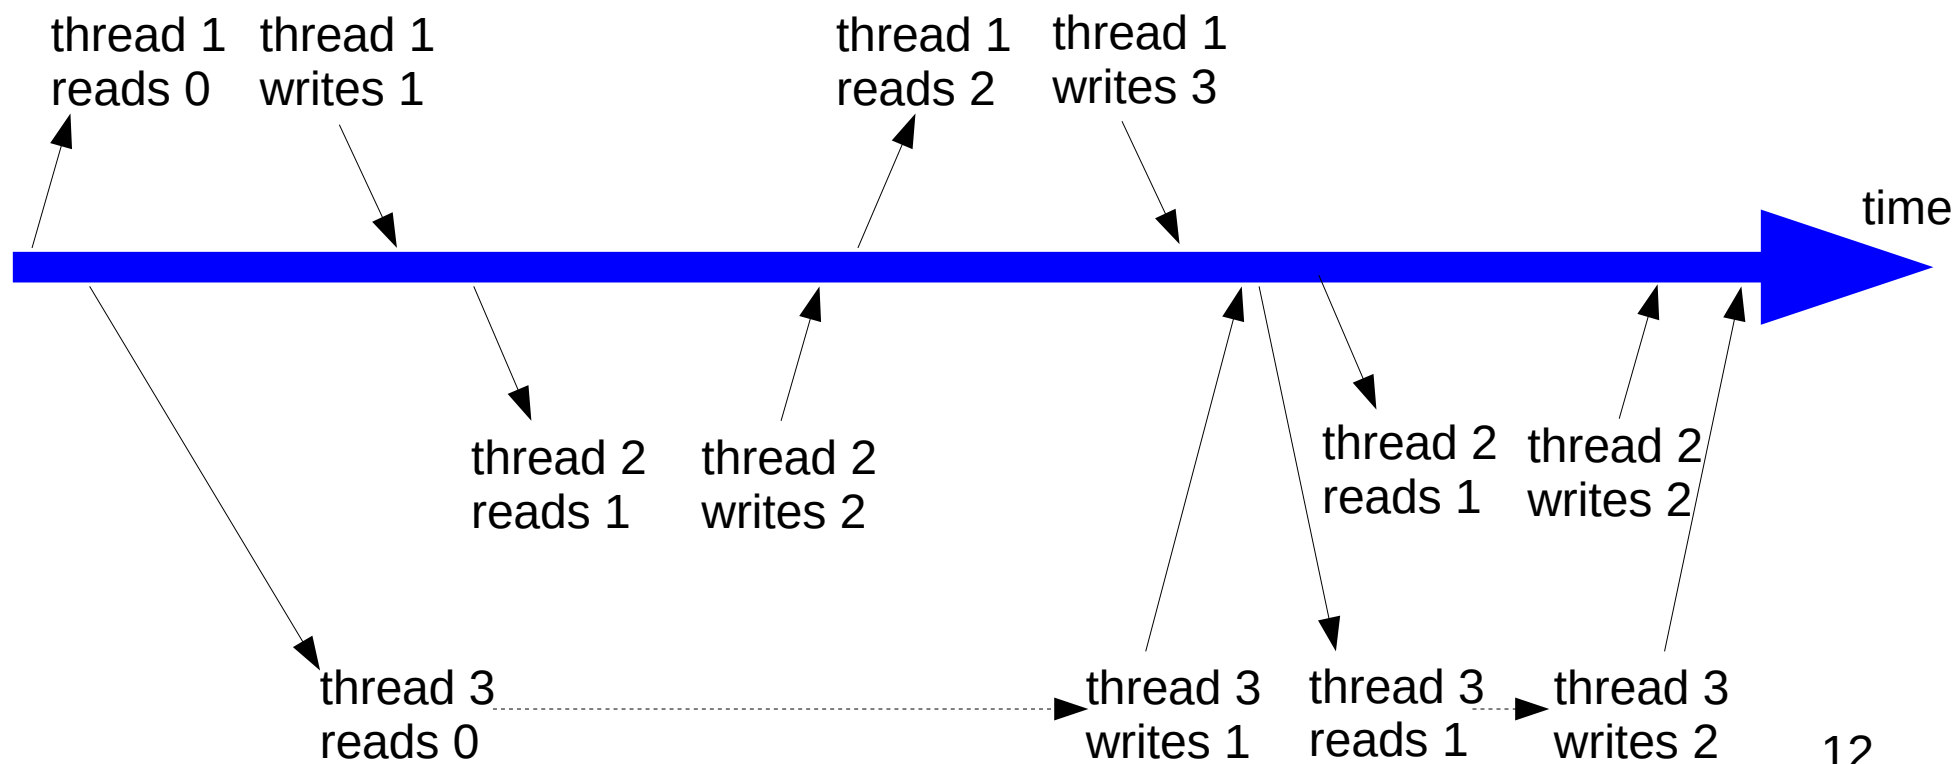
```
volatile int total = 0;
void *incr(void *ptr) {
    for (int i = 0; i < *((int*)ptr); i++)
        total++;
    return NULL;
}

int main() {
    pthread_t tids[NTHREADS];
    int y = NINCR;
    for (int i = 0; i < NTHREADS; i++)
        Pthread_create(&tids[i], NULL, incr, &y);
    for (int i = 0; i < NTHREADS; i++)
        Pthread_join(tids[i], NULL);
    printf("total is: %d\n", total);
    return 0;
}
```

Shared variable

critical section

What Happened?



Threads: Mutual Exclusion

- Need to prevent multiple threads from accessing the same resource at the same time
- In our `sum_array` example, we managed to avoid simultaneous access by giving each thread a separate section of the array
- In general, we'll need a way to temporarily stop a thread while another is accessing the resource it wants to use
 - we use a semaphore or mutex
 - trying to lock a mutex while it is already locked blocks the thread until it is unlocked by the other thread that had already locked it
 - the code between a pair of lock and unlock calls is a *critical section*.

Semaphores

- Special counters with an invariant: their value is never negative
- two *atomic* operations
 - $P(s)$ tries to decrement the counter s (locking the resource), and puts the thread to sleep if the counter is already zero
 - $V(s)$ increments the counter (freeing the resource) and wakes any thread that may be waiting on s
- Mutexes are a subclass of semaphores
 - their value is always either 0 or 1
 - often faster than semaphores because the binary value permits optimizations

Using Semaphores

- Limited resource
 - initialize count to total number of items available
 - decrement just before starting to use one of the items
 - increment when done using the item
- Producer-Consumer
 - initialize count to zero
 - producer generates a new item, then increments semaphore
 - consumer decrements semaphore, then retrieves item

Protecting Shared Resources with a Semaphore

- Suspend execution of thread until resource is “acquired”

```
volatile int total = 0;
sem_t sem;

void *incr(void *ptr) {
    for (int i = 0; i < *ptr; i++) {
        sem_wait(&sem);
        total++;    // CRITICAL SECTION
        sem_post(&sem);
    }
    return NULL;
}

int main() {
    ...
    sem_init(&sem, 0, 1);
```

remember to
initialize the
semaphore first!

Protecting Shared Resources with a Mutex

- Can use a mutex just like a semaphore initialized to 1

```
volatile int total = 0;
pthread_mutex_t M;

void *incr(void *ptr) {
    for (int i = 0; i < *ptr; i++) {
        pthread_mutex_lock(&M);
        total++;      // CRITICAL SECTION
        pthread_mutex_unlock(&M);
    }
    return NULL;
}

int main() {
    ...
    pthread_mutex_init(&M);
```

remember to
initialize the
mutex first!

Problem Solved?

- Sort of...
- Locks in threads are slow.
 - they involve OS calls and *uncached* memory accesses
- Only one instance of the critical section can run at once
 - your eight-core CPU effectively becomes a single-core CPU

Other Solutions

- avoid shared modifiable memory if at all possible
 - (shared read-only memory is OK)
- use a more sophisticated thread synchronization model
 - reader/writer
 - producer-consumer

Problem: Deadlock

- What's about to happen?

Process 1

```
pthread_mutex_lock(&A);  
...  
pthread_mutex_lock(&B);
```

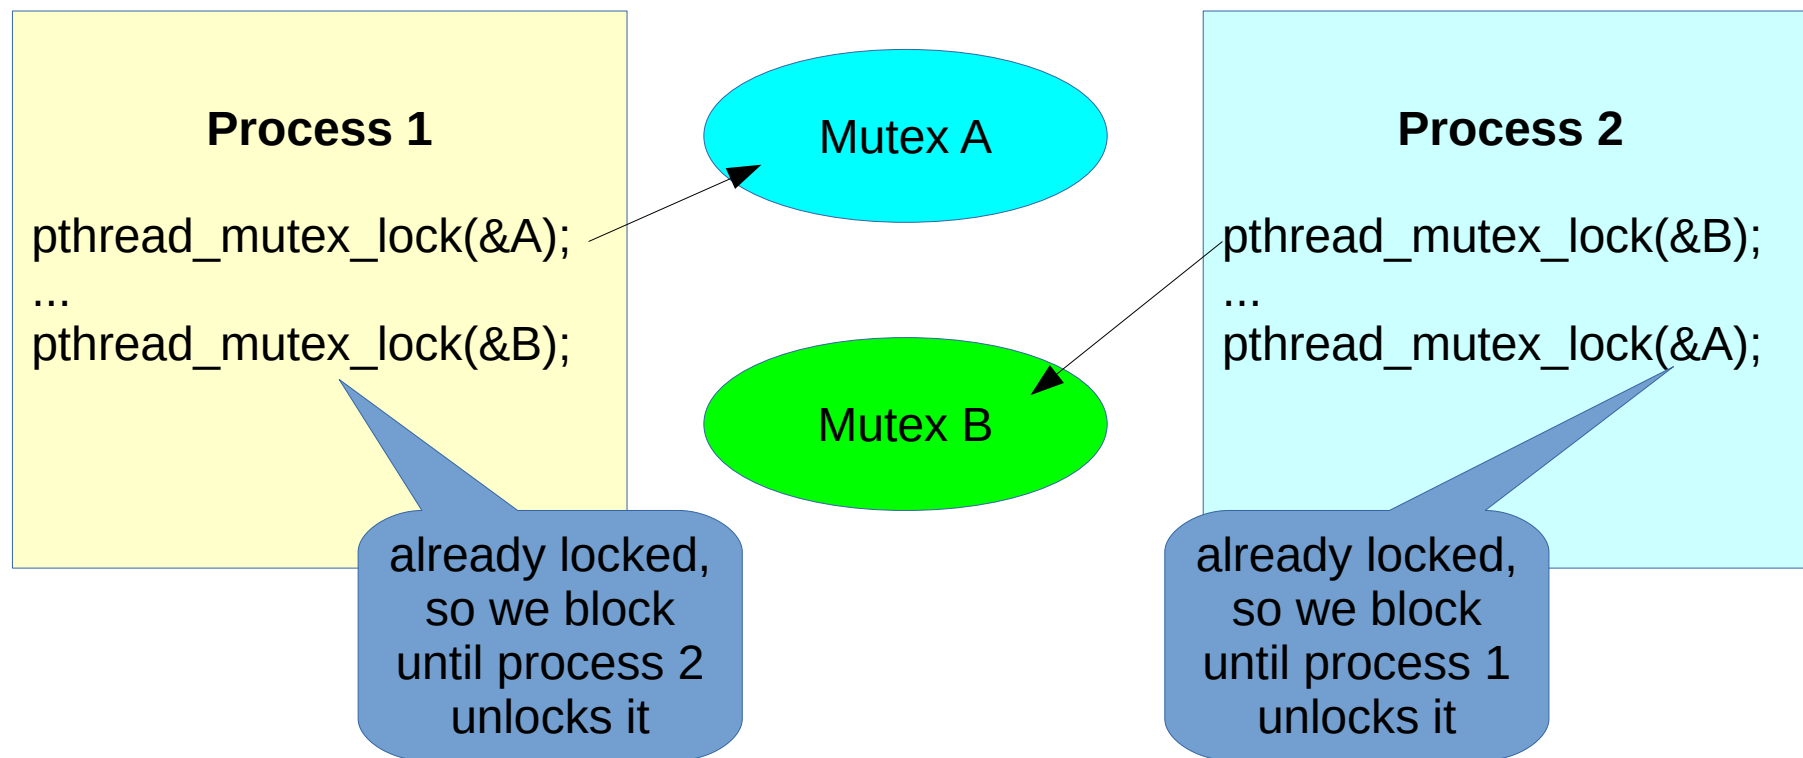
Mutex A

Mutex B

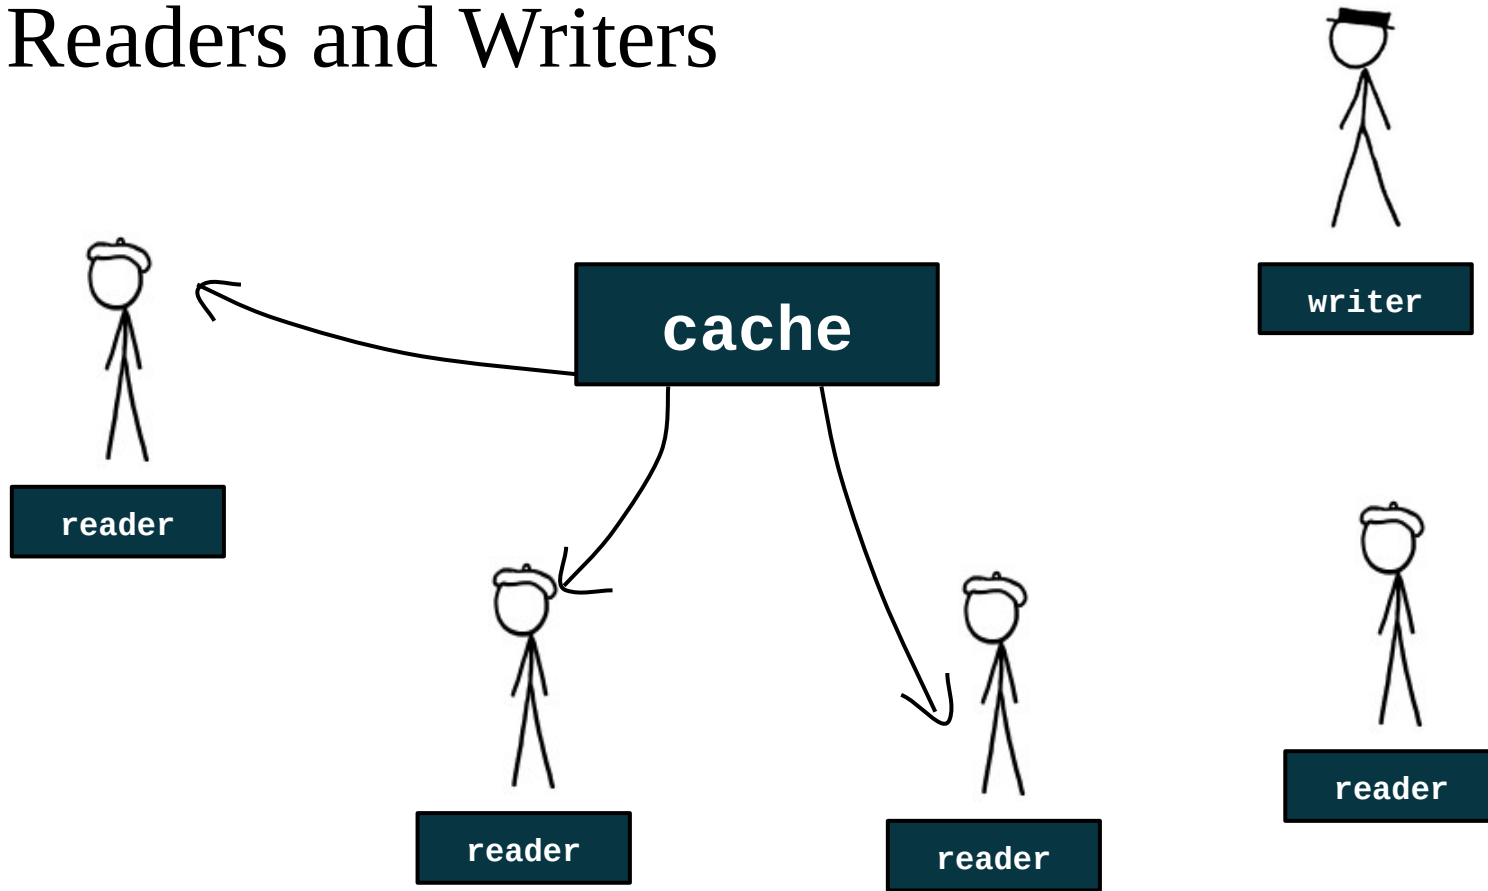
Process 2

```
pthread_mutex_lock(&B);  
...  
pthread_mutex_lock(&A);
```

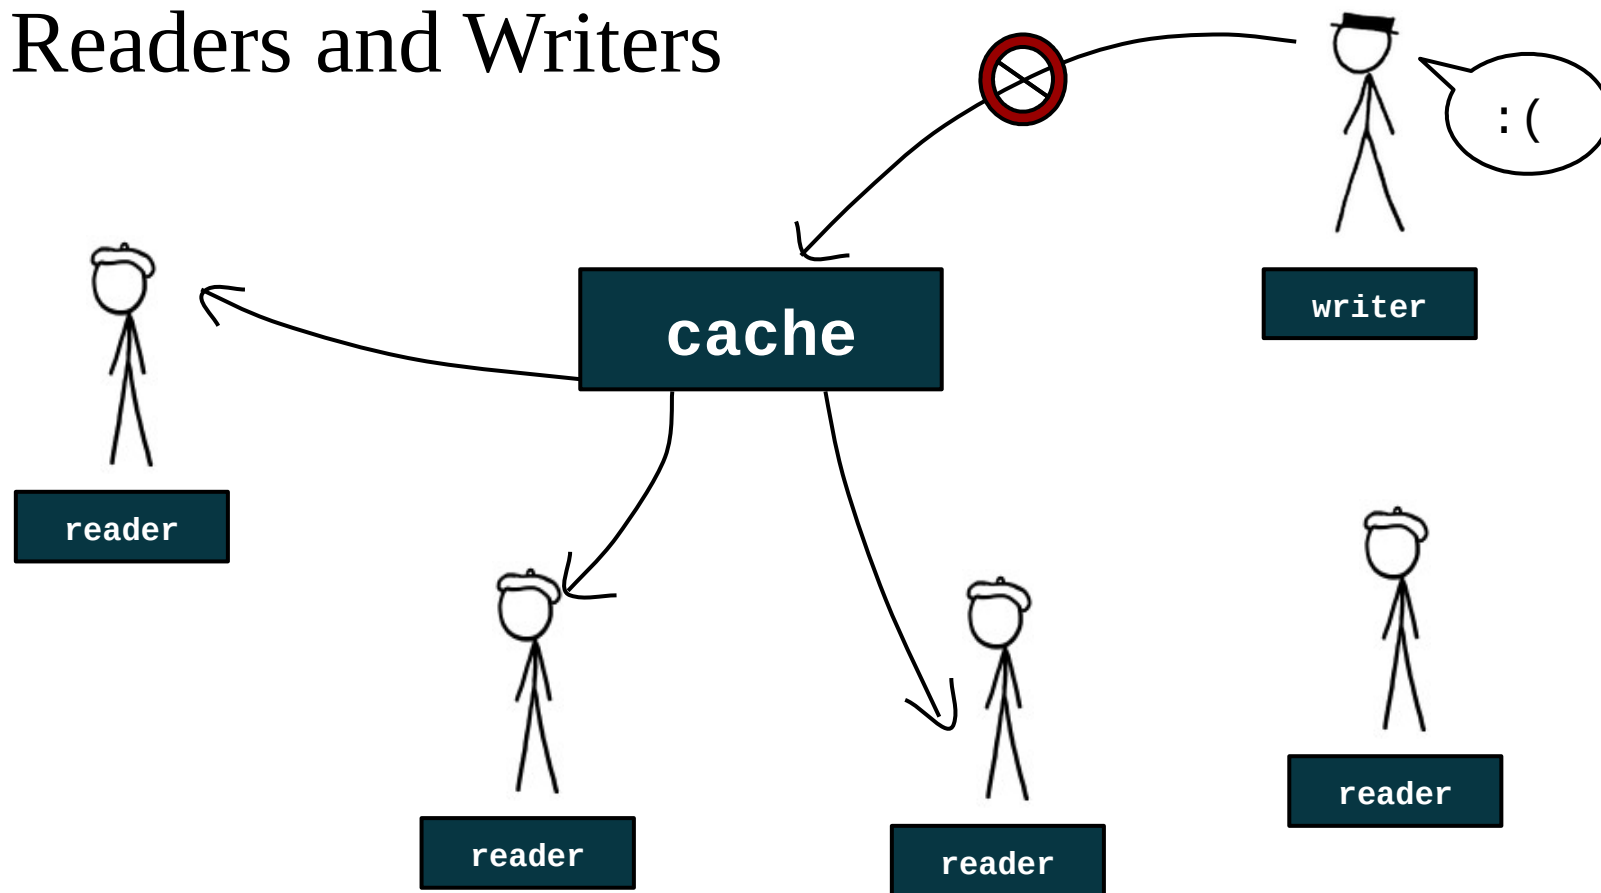
Problem: Deadlock



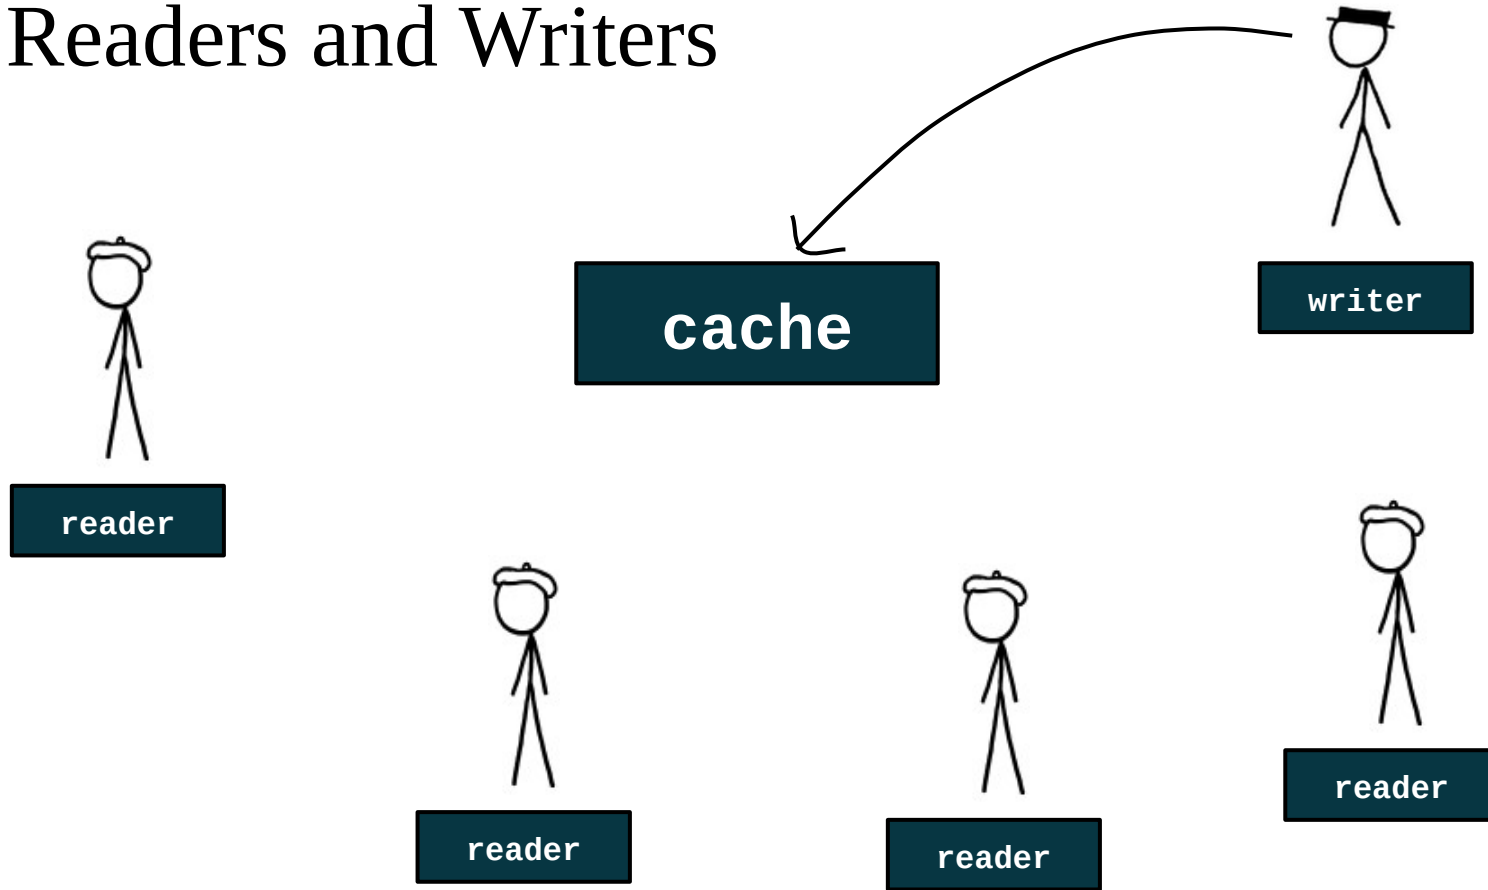
Readers and Writers



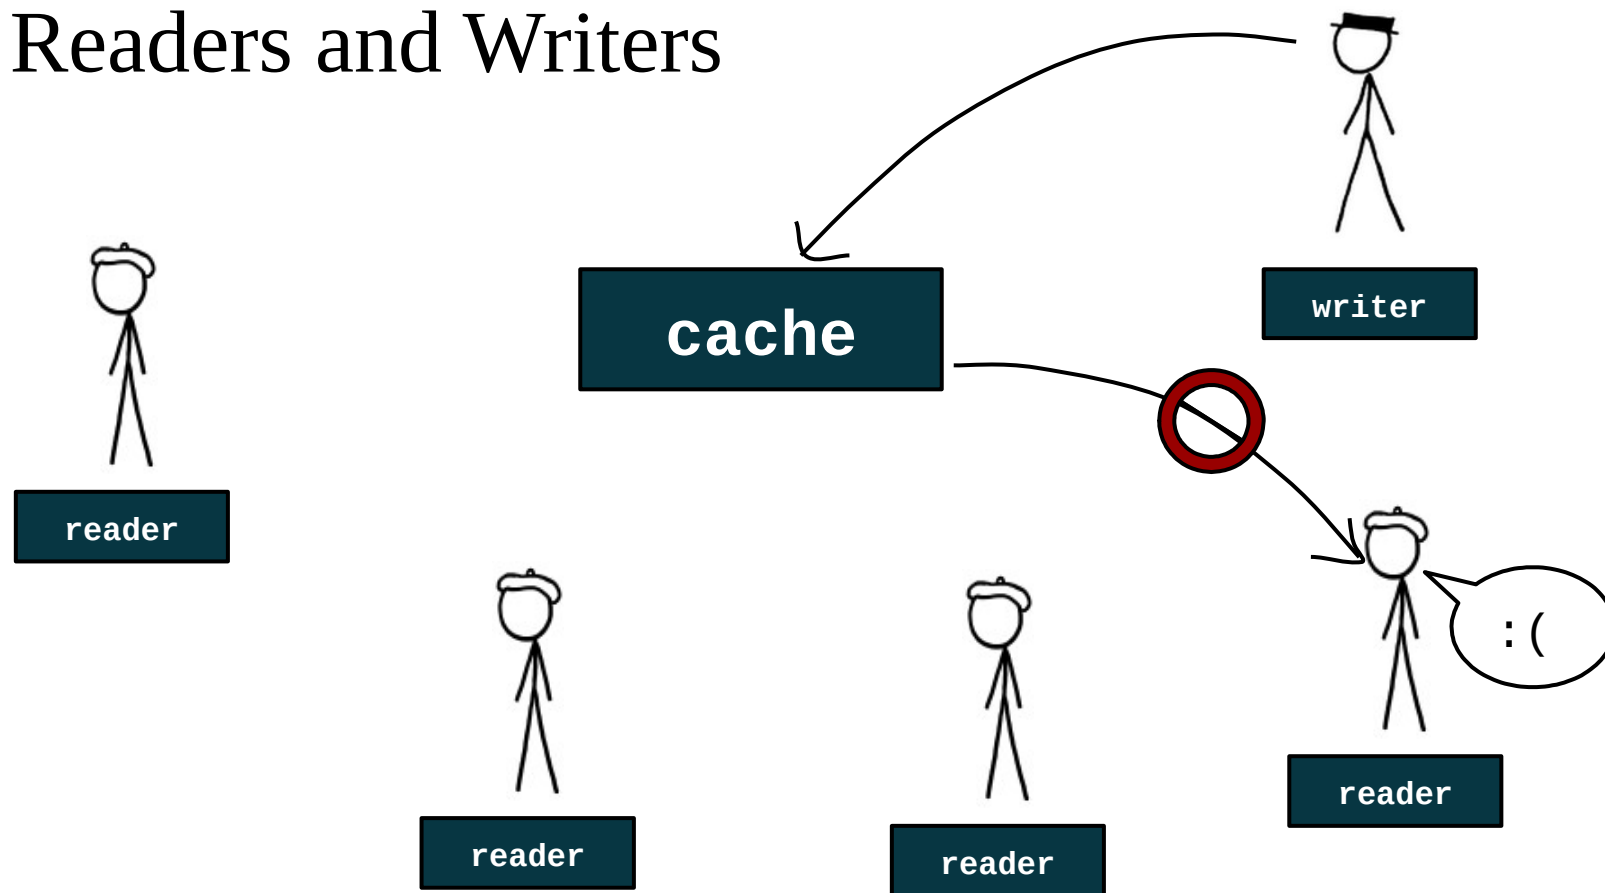
Readers and Writers



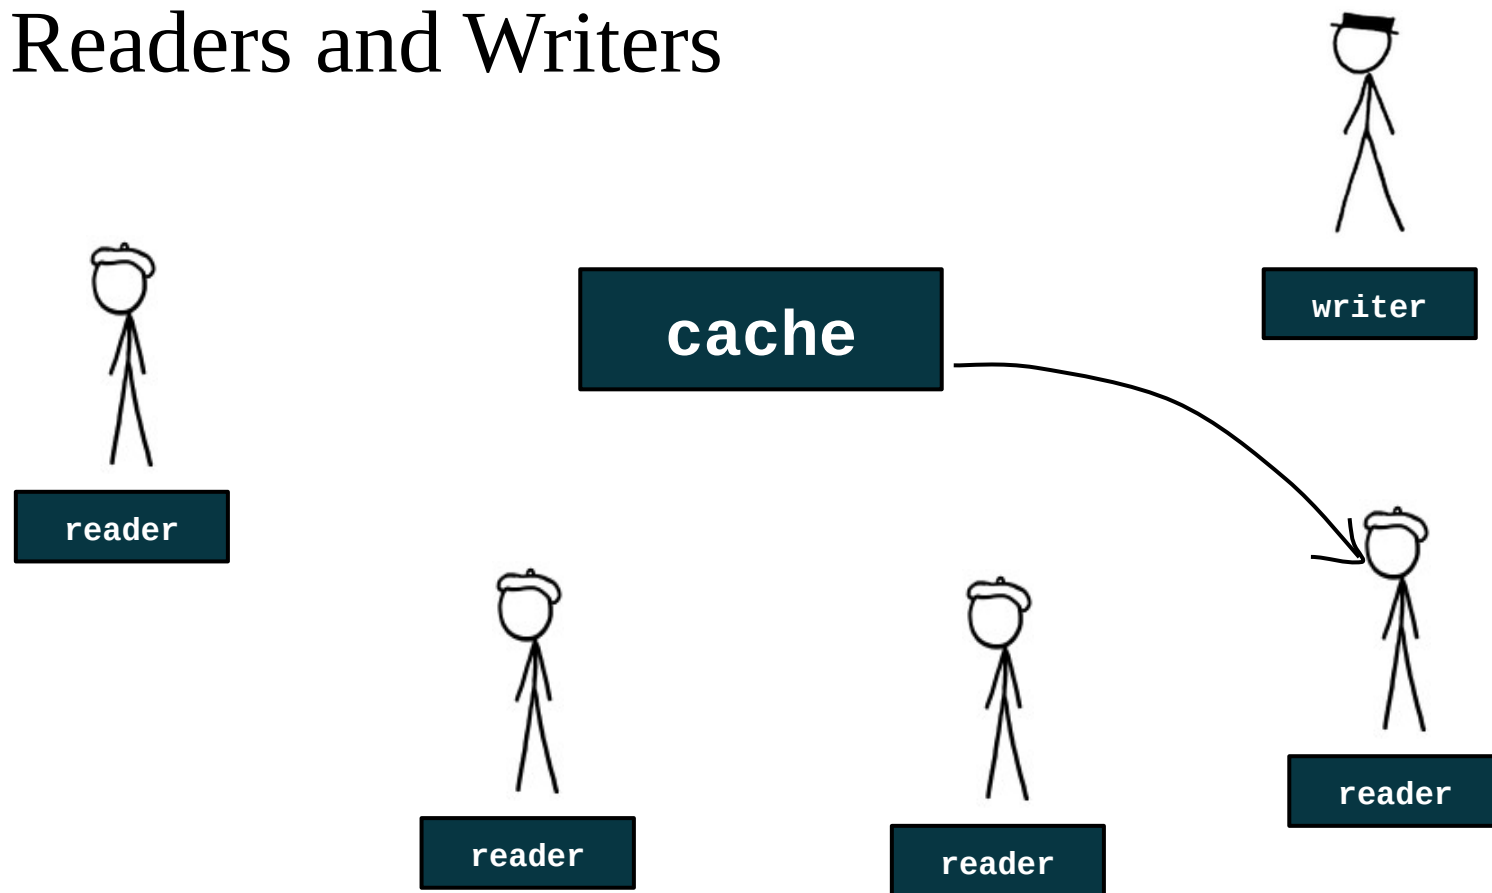
Readers and Writers



Readers and Writers



Readers and Writers



Starvation

- If there are many readers, they may keep writers from acquiring the resource because someone is always reading
 - the writer is being **starved** of the resource
- Important to minimize the amount of time you hold a lock

Read(ers)-Write(r) Lock

- allows a single writer *xor* multiple concurrent readers
- `int pthread_rwlock_init(pthread_rwlock *lock, const pthread_rwlockattr_t *attr);`
- `int pthread_rwlock_rdlock(pthread_rwlock *lock);`
 - lock for reading; blocks if someone is currently (attempting to) write
- `int pthread_rwlock_wrlock(pthread_rwlock *lock);`
 - lock for writing; blocks until all current users finish
- `int pthread_rwlock_unlock(pthread_rwlock *lock);`
- `int pthread_rwlock_destroy(pthread_rwlock *lock);`

Problem: Livelock

- Much like deadlock, but the processes/threads spin indefinitely instead of hanging
- Think of two people trying to get past each other in a hallway
 - both move the same way
 - then both move the other way at the same time
 - awkward dance continues...
- Often the result of attempting to compensate for potential deadlock
 - spinning on a trylock() to avoid hanging on a lock()

Which Lock Do I Use?

- Consider **what** is shared and what type of **access** is desired
 - only one thread at a time allowed
 - ?
 - more than one instance of a shared resource is available
 - ?
 - multiple threads can read concurrently, but only one may write at a time
 - ?

Which Lock Do I Use?

- Consider **what** is shared and what type of **access** is desired
 - only one thread at a time allowed
 - **mutex** example: global count variable
 - more than one instance of a shared resource is available
 - **semaphore** example: multiple free slots in a shared buffer
 - multiple threads can read concurrently, but only one may write at a time
 - **readers-writer lock** example: lookup from global list, web-proxy cache